

# Scientific Computing for Biologists

Biology 723

Fall 2012

Tue 2:50-5:20

Instructor: Paul M. Magwene

Email: [paul.magwene@duke.edu](mailto:paul.magwene@duke.edu)

Phone: 613-8159

28 August 2012

# Overview of Lecture

- Course Mechanics

- Goals of course
- Structure of lectures
- Grading

- Introduction to R

- Advantages of R and Python
- R and Python Resources
- Important programming concepts
- Introduction to data types and data structures in R and Python
- Literate programming

- Hands-On Session

# Class Structure

## ■ Lectures

- Typically 60-75 minutes
- Emphasize the mathematical basis of the methods/approaches from both a geometric and algebraic basis
- Discuss algorithms underlying the methods

## ■ Hands-on

- Walk through some examples
- Apply the techniques and concepts to real data
- Highlight available R/Python libraries

# Syllabus

<i>Date</i>	<i>Topic</i>
August 28	Introduction; Getting Acquainted with R
September 4	Data as vectors; Vector operations: dot product, correlation, regression as projection
September 11	Linear Algebra Review I; Descriptive statistics as matrix operations, multivariate visualizations
September 18	Linear Algebra Review II; Regression models
September 25	Eigenvectors and Eigenvalues; Principal Components Analysis
October 2	Singular Value Decomposition, Biplots, and Correspondence Analysis
October 9	Discriminant analysis and Canonical Variate Analysis
October 16	Fall Break
October 23	Analyses based on Similarity/Distance I; Hierarchical and K-means clustering
October 30	Analyses based on Similarity/Distance II; Multidimensional scaling
November 6	Randomization and Monte Carlo Methods; Jackknife, Bootstrap, and Permutation Tests
November 13	Building Simulations in Python
November 20	Building Bioinformatics Pipelines I; Pipes, redirection, subprocesses
November 27	Building Bioinformatics Pipelines II; Putting the concepts to work

# Course Goals

- 1 Introduce multivariate statistics from a geometric perspective, emphasizing the geometry of vector spaces.
- 2 Develop a good working knowledge of R (a statistical computing environment) and Python (a general purpose programming language).
- 3 Provide the tools and knowledge to conduct reproducible computational and statistical research.

# Texts for Course

- Wickens, T. D. 1995. The geometry of multivariate statistics.
- Matloff, N. 2011. The Art of R Programming.
- Downey, A. B., J. Elkner, and C. Meyers. How to think like a computer scientist: learning with Python.
  - Available at <http://www.ibiblio.org/obp/thinkCSpy/>

# Supplementary Texts

## ■ Statistics

- Krzanowski, W. J. 2003. Principles of multivariate analysis. Oxford University Press.
- Sokal, R. R. and F. J. Rohlf. 1995. Biometry. W. H. Freeman.

## ■ Math

- Hamilton, A. G. 1989. Linear algebra: an introduction with concurrent examples. Cambridge University Press.

# Grading

- Problem sets/programming assignments
  - 10 homeworks over the course of the semester
  - Programming assignments should not be submitted until they produce correct results; I will provide you with scripts/corresponding data to check the correctness of your code
  - No credit for late assignments
  - Overall grade: assignments are worth one point each, for a total of 10 points for the semester. Final grades will be assigned based on this scale: 9+ pts = A, 8 = B, 7 = B-, 6 = C, 5 = F



# Why Both R and Python?

The first half of the course is built around these use of the R programming language; in the second half of the course we will use both R and Python.

- R is geared toward statistical computing
  - Great set of built-in facilities for statistically oriented tasks
  - Somewhat cumbersome syntax for non-statistical tasks
- Python is a general programming language
  - Clearer syntax
  - Wider range of modules
    - web programming, databases, numerical analysis, etc.
  - More natural language for simulation
  - More suitable as a 'glue' language
    - building bioinformatics pipelines

# Introduction to R

# What is R?

- 'A language and environment for statistical computing and graphics'
- First developed in the mid-90s
- Derives from the S language
  - S was developed at Bell Labs in the mid-80s
- Advantages
  - Free and open-source
  - Much of the academic statistical community has adopted it
  - Active developer and user community
  - Wealth of built-in and user contributed libraries available for all types of analyses
- Disadvantages
  - GUI not as well developed as commercial statistical packages
    - S-Plus; site licensed by Duke - see OIT website
  - Has higher learning curve than some other simpler statistical software
  - Command-line can be intimidating

# R Resources on the Web

- Home Page
  - <http://www.r-project.org>
- Comprehensive R Archive Network (CRAN)
  - <http://cran.r-project.org/mirrors.html>
  - See especially the 'Task Views'
    - Statistical and population genetics
    - Environmental and ecological analysis
    - Spatial statistics
- Introductions and Tutorials
  - see <http://cran.r-project.org/other-docs.html>

## Some R Packages of Interest

- Bioconductor – software package geared towards analysis of genomic data, especially microarray data,  
<http://www.bioconductor.org/>
- ape – ‘Analysis of Phylogenetics and Evolution’,  
<http://ape.mp1.ird.fr/>
- ade4 – Analysis of Ecological Data : Exploratory and Euclidean methods in Environmental sciences,  
<http://pbil.univ-lyon1.fr/ADE-4/home.php?lang=eng>

# Some Important Programming Concepts

## ■ Data Types

- refer to the types of values that can be represented in a computer program
- determine the representation of values in memory
- determine the operations you can perform on those values
- Examples: integers, strings, floating point values

## ■ Data Structures

- a way of storing collections of data
- different structures are more efficient for particular types of operations
- Examples: lists, hash tables, stacks, queues, trees

## ■ Variables

- Variables are references to objects/values in memory
- Think of them as labels that point to particular places in a computer's memory

# More Important Programming Concepts

## ■ Statement

- an instruction that a computer program can execute
- Example: `print("Hello, World!")`

## ■ Operators

- Symbols representing specific computations
- Example: `+`, `-`, `*` (addition, subtraction, multiplication)

## ■ Expression

- a combination of values, variables, and operators
- Example: `1 + 1`

## ■ Functions (subroutines, procedures, methods)

- A piece of code that carries out a specific task, set of instructions, calculations, etc.
- Typically used to encapsulate algorithms

# Basic Data Types, Data Structures and Operators in R



# Numeric Data Types in R

## ■ Floating point values ('doubles')

```
> x <- 10.0  
> typeof(x)  
[1] "double"
```

## ■ Complex numbers

```
> x <- 1+1i  
> typeof(x)  
[1] "complex"
```

## ■ Integers

- Default numeric type is double, must explicitly ask for integers if single values

```
> x <- as.integer(10)  
> typeof(x)  
[1] "integer"
```

# Additional Data Types in R

## ■ Boolean('logical')

```
> x <- TRUE # or x <- T
> x <- F # or x <- FALSE
> typeof(x)
[1] "logical"
```

## ■ Character strings

```
> x <- 'Hello' # or x <- "Hello"
> typeof(x)
[1] "character"
```

# Arithmetic Operators and Mathematical Functions in R

```
> 10 + 2 # addition
[1] 12
> 10 - 2 # subtraction
[1] 8
> 10 * 2 # multiplication
[1] 20
> 10 / 2 # division
[1] 5
> 10 ^ 2 # exponentiation
[1] 100
> 10 ** 2 # alternate exponentiation
[1] 100
> sqrt(10) # square root
[1] 3.162278
> 10 ^ 0.5 # same as square root
[1] 3.162278
> pi*(3)**2 # R knows some useful constants
[1] 28.27433
> exp(1) # exponential function
[1] 2.718282
```

# Simple Data Structures in R: Vectors

Vectors are the simplest data structure in R

- vectors represent an ordered list of items

```
> x <- c(2,4,6,8)
> y <- c('joe', 'bob', 'fred')
```

- vectors have length (possibly zero) and type

```
> typeof(x)
[1] "double"
> length(x)
[1] 4
> typeof(y)
[1] "character"
```

# Simple Data Structures in R: Vectors

Accessing the objects in a vector is accomplished by 'indexing':

- The elements of the vector are assigned indices  $1 \dots n$  where  $n$  is the length of the vector

```
> x <- c(2,4,6,8)
```

```
> length(x)
```

```
[1] 4
```

```
> x[1]
```

```
[1] 2
```

```
> x[2]
```

```
[1] 4
```

```
> x[3]
```

```
[1] 6
```

```
> x[4]
```

```
[1] 8
```

# Simple Data Structures in R: Vectors

- Single objects are usually represented by vectors as well

```
> x <- 10.0  
> length(x)  
[1] 1  
> x[1]  
[1] 10
```

- Every element in a vector is of the same type

- If this is not the case the the values are coerced to enforce this rule

```
> x <- c(1+1i, 2+1i, 'Fred', 10)  
> x  
[1] "1+1i" "2+1i" "Fred" "10"
```

# Arithmetic Operators Work on Vectors in R

Most arithmetic operators work element-by-element on vectors in R

```
> x <- c(2, 4, 6, 8)
> y <- c(0, 1, 2, 3)
> x + y
[1] 2 5 8 11
> x - y
[1] 2 3 4 5
> x * y
[1] 0 4 12 24
> x^2
[1] 4 16 36 64
> sqrt(x)
[1] 1.414214 2.000000 2.449490 2.828427
```

# Simple Data Structures in R: Lists

## Lists

- Lists in R are like vectors but the elements of a list are arbitrary objects (even other lists)

```
> x <- list('Bob', 27, 10, c(720, 710))
```

```
> x
```

```
[[1]]
```

```
[1] "Bob"
```

```
[[2]]
```

```
[1] 27
```

```
[[3]]
```

```
[1] 10
```

```
[[4]]
```

```
[1] 720 710
```



# Simple Data Structures in R: Lists

## Accessing objects in Lists:

- Items in lists are accessed in a different manner than vectors.
  - Typically you use double brackets (`[[ ]]`) to return the element at index `i`
  - Single brackets always return a list containing the element at index `i`

```
> x <- list('Bob', 27, 10, c(720,710))  
> typeof(x[1])  
[1] "list"  
> typeof(x[[1]])  
[1] "character"
```

# Simple Data Structures in R: Lists

## ■ Objects in R lists can be named

```
> x <- list(name='Bob', age=27, years.in.school=10)
> x
$name
[1] "Bob"

$age
[1] 27

$years.in.school
[1] 10
```

## ■ Named list objects can be accessed via the \$ operator

```
> x$years.in.school
[1] 10
> x$name
[1] "Bob"
```

## ■ The names of list objects can be accessed with the names() function

```
> names(x)
[1] "name" "age" "years.in.school"
```

# Things to Remember

- Try it out - programming involves experimentation
- Practice - learning to program, like learning a foreign language, requires lots of practice.
- Persist - many new tools/concepts can be hard to grasp at first. Keep plugging away until you get that 'Aha!' moment

## You might be surprised to find that...

- Programming is fun! (at least sometimes)
- Math is fun! (at least sometimes)
- Statistics is fun! (at least sometimes)
- Gaining new insights into how your biological system of interest works is fun! (always)